

Top 50 ADF Real-Time Scenarios

(Full Explanation + How-To + Example)



Aditya Raja

<https://www.linkedin.com/in/adityaraja853/>

Top 50 ADF Real-Time Scenarios (Full Explanation + How-To + Example)

◆ Scenario 1 — Load Only New/Modified Records (Incremental Load / Delta Load)

📌 Explanation

Most pipelines should not reload full tables. Use a **watermark** column (`LastUpdated, ModifiedDate`) to load only changed rows.

📌 How To Implement

1. Create Watermark table.
2. Lookup last watermark value in ADF.
3. Use dynamic SQL in Copy Activity:
4. `SELECT * FROM Sales`
5. `WHERE ModifiedDate > @{activity('GetWatermark').output.firstRow.lastValue}`
6. Run Copy → update watermark using Stored Proc.

📌 Real Example

Last watermark = **2025-01-10 12:00**

Pipeline loads only rows newer than that value.

◆ Scenario 2 — Trigger Pipeline Only When File Exists

📌 Explanation

Avoid failures when file hasn't arrived yet.

📌 How To Implement

1. Use **Get Metadata** → check `ChildItems`.
2. Use **If Condition**:
 - `TRUE` → Copy Activity
 - `FALSE` → Log + Skip / Send alert

📌 Real Example

If `orders_20250114.csv` exists → process
Else → send Teams/Email alert.

◆ Scenario 3 — Trigger Pipeline When File Arrives (Event-Based Trigger)

📌 Explanation

ADF supports automatic trigger using **BlobCreated** events.

📌 How To Implement

1. Create Event Trigger.
2. Select Storage Account + Container.
3. Add filter path `/incoming/orders`.

📌 Example

When a new file arrives in `adls/raw/orders/`, pipeline runs instantly.

◆ Scenario 4 — Pass Parameters from ADF to Databricks Notebook

📌 Explanation

Dynamic pipelines need to instruct Databricks which file/date/table to process.

📌 How To Implement

ADF Notebook Activity:

```
baseParameters:  
  fileDate: @pipeline().parameters.fileDate  
  objectName: @pipeline().parameters.tableName
```

Inside notebook:

```
fileDate = dbutils.widgets.get("fileDate")  
table = dbutils.widgets.get("objectName")
```

📌 Example

fileDate = 2025-01-14

Databricks loads only that partition from ADLS.

◆ Scenario 5 — Implement SCD Type-2 (Maintain Data History)

📌 Explanation

SCD2 maintains full history when dimension data changes.

📌 How To Implement

In Mapping Data Flow:

- Use **Alter Row**:
 - INSERT when new
 - UPDATE when data changed
- Add:
 - EffectiveStartDate
 - EffectiveEndDate
 - IsCurrent flag

📌 Example

Customer moved from “Delhi → Bangalore”

Old record: EndDate updated

New record inserted with IsCurrent = 1.

◆ Scenario 6 — Merge Incremental Changes Using Databricks (Delta MERGE)

📌 Explanation

Databricks Delta supports ACID merges for incremental updates.

📌 How To Implement

```
MERGE INTO silver.customers AS tgt
USING bronze.customers AS src
ON tgt.id = src.id
WHEN MATCHED THEN UPDATE SET *
WHEN NOT MATCHED THEN INSERT *
```

📌 Example

Used when daily updates need to sync with Silver layer.

◆ Scenario 7 — Handle Schema Drift Automatically

📌 Explanation

New columns appear frequently in JSON/CSV files.

📌 How To Implement

In **Mapping Data Flow**:

- Enable Schema Drift
- Use Auto-Map
- Derive missing columns if needed

📌 Example

Incoming JSON adds new field `device_type` → pipeline continues without failing.

◆ Scenario 8 — Load Nested JSON Files

📌 Explanation

ADF Mapping Dataflow supports flattening nested JSON.

📌 How To Implement

1. Use **Flatten Transformation**.
2. Select nested arrays/objects.
3. Map flattened fields to sink.

📌 Example

JSON:

```
customer.address.city
```

Flatten → Create column: `city`.

◆ Scenario 9 — Optimize Copy Activity for Large Data (10M+ Records)

📌 Explanation

To speed up SQL → ADLS or ADLS → Synapse loads.

📌 How To Implement

- Enable **parallel copy**
- Increase **DIUs**
- Use **partition option**
- Use **staging** for SQL → Synapse

📌 Example

Copy 20M rows from SQL → ADLS in 12 minutes using 16 parallel copies.

◆ Scenario 10 — Process Multiple Files in Folder Dynamically

📌 Explanation

ADF must ingest multiple files automatically.

📌 How To Implement

1. Get Metadata → ChildItems
2. ForEach → loop through file list
3. Use dynamic file path in Copy:
4. `@item().name`

📌 Example

Folder contains:

- sales_20250101.csv
- sales_20250102.csv
- sales_20250103.csv

Pipeline loads all files one by one.

◆ Scenario 11 — Process millions of rows efficiently (SQL → ADLS / Synapse)

📌 Scenario

You need to move **tens of millions of rows** daily from SQL to ADLS/Synapse and the pipeline is very slow.

📌 Explanation

For large volumes, **parallelism + partitioning + proper sink** (Parquet/PolyBase) are key. A single-threaded copy or row-by-row insert will be too slow.

📌 How to handle

1. Use **PolyBase / Bulk insert** when loading into Synapse.
2. **Partition the data** (date, id ranges) and copy per partition.
3. Use **Copy Activity** with:
 - o maxConcurrentConnections
 - o parallelCopies
4. Prefer **Parquet** or **Delta** instead of CSV.
5. Use **staging** (e.g., Blob → Synapse via PolyBase).

📌 Example

- In Copy Activity → **Source query**:

```
SELECT * FROM Sales
WHERE OrderDate >= @startDate
AND OrderDate < @endDate;
```

- In pipeline:
 - o Use a **ForEach** over date partitions (e.g., last 7 days).
 - o Each iteration calls the Copy Activity with different startDate/endDate.

◆ Scenario 12 — Dynamic file-name based ingestion (daily files)

📌 Scenario

Every day, a new file like `sales_2025-11-23.csv` lands in ADLS. You must build **one reusable pipeline** that picks today's file.

📌 Explanation

You can use **pipeline parameters + dynamic content** to avoid hardcoding file names.

📌 How to handle

1. Create a **pipeline parameter**: `fileDate`.
2. Default value: `@formatDateTime(utcNow(), 'yyyy-MM-dd')`

3. In your **dataset file path**, use:

```
@concat('sales_', pipeline().parameters.fileDate, '.csv')
```

4. Trigger the pipeline **daily** with a schedule trigger.

📌 Example

- Dataset → File name:
`@concat('sales_', pipeline().parameters.fileDate, '.csv')`
- If you need to re-run for a specific day, manually trigger and **override fileDate** (e.g. 2025-11-20).

◆ Scenario 13 — Call REST API with pagination (multiple pages of data)

📌 Scenario

A REST API returns only **1000 records per page** and uses a `nextPageToken` for pagination. You must **load all pages** into ADLS or SQL.

📌 Explanation

You use **Web Activity + Until loop** to repeatedly call the API until there's no next page.

📌 How to handle

1. Start with an initial Web Activity to call Page 1.
2. Store `nextPageToken` in a **pipeline variable**.
3. Use an **Until** activity:
 - Condition: `@equals(variables('nextPageToken'), null)`
4. Inside **Until**:
 - Web Activity calls API with `nextPageToken`.
 - Append response to ADLS via Copy Activity (REST → ADLS).
 - Update `nextPageToken` variable from response.

📌 Example

- URL in Web Activity:

```
@concat('https://api.example.com/users?pageToken=', variables('nextPageToken'))
```

- Variable update:

```
setVariable:  
Name: nextPageToken  
Value: @activity('CallAPI').output.nextPageToken
```

◆ Scenario 14 — Process all files in a folder (loop through multiple files)

📌 Scenario

A folder `/raw/sales/` contains **multiple daily files**. You must process **all files automatically**.

📌 Explanation

Use **Get Metadata** to list files and then **ForEach** to loop through each file name.

📌 How to handle

1. **Get Metadata** activity:
 - o Point to folder `/raw/sales/`
 - o Set **Field list** = `Child items`
2. **ForEach** activity:
 - o Items: `@activity('Get Metadata1').output.childItems`
3. Inside ForEach:
 - o Copy Activity with dataset file path:
 - Directory: `/raw/sales/`
 - File name: `@item().name`

📌 Example

Dataset file field:

```
@item().name
```

This will copy/process each file like `sales_2025-11-20.csv`, `sales_2025-11-21.csv`, etc, without hardcoding.

◆ Scenario 15 — Prevent parallel pipeline runs (no overlap allowed)

📌 Scenario

Your daily pipeline should **not run if the previous run is still running** (to avoid double-processing).

📌 Explanation

Use **pipeline concurrency** setting so that ADF **queues** new runs instead of running them in parallel.

📌 How to handle

1. Go to **pipeline Settings**.
2. Set **Concurrency** = 1.
3. If a new trigger fires while pipeline is running, the new run is **queued**, not executed in parallel.

📌 Example

- Daily trigger at 1 AM.
 - Suppose 1 AM run takes 2 hours.
 - If you manually trigger the same pipeline at 1:30 AM, the second run will wait until the first finishes.
-

◆ Scenario 16 — Audit logging for every run (who, when, status)

📌 Scenario

You need an **audit log** table capturing: pipeline name, runId, startTime, endTime, status, errorMessage.

📌 Explanation

Use a **Stored Procedure activity** or **Copy to log table** at the end (or on failure) of the pipeline.

📌 How to handle

1. Create a SQL table PipelineAuditLog.
2. At the end of the pipeline, use **Stored Procedure** activity:
 - Pass:
 - @pipeline().RunId
 - @pipeline().Pipeline (name)
 - @utcNow() for end time
 - activity('YourActivity').Status or full output.
3. For failures:
 - Use an extra activity in **On Failure** path to log error details.

📌 Example

Stored procedure:

```
CREATE PROCEDURE usp_InsertPipelineLog
(
    @RunId NVARCHAR(100),
    @PipelineName NVARCHAR(200),
    @Status NVARCHAR(50),
    @StartTime DATETIME2,
    @EndTime DATETIME2,
    @ErrorMessage NVARCHAR(MAX)
)
AS
BEGIN
    INSERT INTO PipelineAuditLog
    VALUES (@RunId, @PipelineName, @Status, @StartTime, @EndTime, @ErrorMessage);
END
```

◆ Scenario 17 — Validate file format / schema before processing

📌 Scenario

If the **file has wrong number of columns** or **wrong extension**, the pipeline should **fail early** and not load junk.

📌 Explanation

Use **Get Metadata** + **If Condition** and optionally **Data Flow Assert** / schema validation.

📌 How to handle

1. **Get Metadata** on file:
 - o Fields: `columnCount` (for some sources), `itemName`, `size`.
2. **If Condition**:
 - o Example condition: `@equals(endsWith(activity('Get Metadata').output.itemName, '.csv'), true)`
3. Inside True (valid): run Copy/Dataflow.
4. Inside False: call **Fail Activity** or log error.

📌 Example

- Expression to check extension:

```
@equals(last(split(activity('Get Metadata1').output.itemName, '.')), 'csv')
```

If not `.csv`, pipeline goes to error handling path.

◆ Scenario 18 — Process large JSON files via Databricks (or Mapping Data Flow)

📌 Scenario

You receive **very large JSON files** with nested structure. Mapping Data Flow struggles or is too expensive; you prefer Databricks.

📌 Explanation

Databricks **Auto Loader** or Spark can handle large JSON with **schema inference + incremental load**.

📌 How to handle (Databricks)

1. **ADF Notebook Activity** triggers Databricks notebook.
2. In notebook:

```
df = spark.read.json("abfss://raw@youraccount.dfs.core.windows.net/json/")
# Flatten, transform, filter
df_flat = df.select("id", "name", "address.city", "address.country")
df_flat.write.format("delta").mode("append").save("/mnt/silver/customers")
```

3. ADF orchestrates this as part of pipeline, before loading to Synapse/Gold.

📌 Example

- ADF pipeline:
 - Copy file from landing → raw.
 - Notebook Activity to process and write into **Delta tables**.
-

◆ Scenario 19 — Copy from on-prem Oracle/SQL Server to ADLS using Self-Hosted IR

📌 Scenario

Your data source is **on-prem SQL/Oracle**, not accessible over public internet. You must securely copy data to ADLS.

📌 Explanation

You need a **Self-Hosted Integration Runtime (SHIR)** installed inside the network/VPN.

📌 How to handle

1. Install **Self-Hosted IR** on a VM or on-prem server with DB access.
2. In ADF, create **Linked Service (SQL/Oracle)** using SHIR.
3. Use Copy Activity:
 - Source: on-prem SQL/Oracle (linked with SHIR).
 - Sink: ADLS Gen2.
4. Use **SQL query** or **table** selection.

📌 Example

Source dataset uses Linked Service `OnPremSql_LS` with IR type = `SelfHostedIntegrationRuntime1`.

SQL query:

```
SELECT * FROM Orders WHERE OrderDate >= @lastWatermark;
```

◆ Scenario 20 — Stored procedure driven ETL (using SQL logic inside ADF pipeline)

📌 Scenario

You want the **business/merge logic to stay in SQL (SP)** while ADF handles orchestration.

📌 Explanation

ADF can call **Stored Procedures** as part of the pipeline to execute complex logic like SCD2, temp tables, cleanup, etc.

📌 How to handle

1. Create a **Stored Procedure activity**.
2. Linked Service → your SQL DB.
3. Pass parameters (e.g., RunDate, Watermark) from pipeline.
4. SP does:
 - o staging → target table merge
 - o update watermark table
 - o log row counts.

📌 Example

- In ADF Stored Procedure activity parameters:

```
RunDate = @utcNow()  
Watermark = @pipeline().parameters.watermark
```

- Sample SP pseudocode:

```
CREATE PROCEDURE usp_ETL_Sales (@RunDate DATETIME2, @Watermark DATETIME2)  
AS  
BEGIN  
    -- Load from staging to final table with MERGE  
    MERGE dbo.SalesFinal AS T  
    USING dbo.SalesStaging AS S  
    ON T.Id = S.Id  
    WHEN MATCHED THEN UPDATE SET ...  
    WHEN NOT MATCHED THEN INSERT (...);  
  
    UPDATE ControlTable SET LastWatermark = @RunDate WHERE ProcessName='Sales';  
END
```

◆ Scenario 21 — Partition data in Mapping Data Flows for performance

📌 Explanation (What & Why)

When you process large datasets in Mapping Dataflows, a single partition can become a bottleneck. Partitioning lets ADF split the data (e.g., by date or ID) so Spark can process partitions in parallel → faster pipelines and better cluster utilization.

📌 How to implement in ADF (Steps)

1. Open your **Mapping Data Flow**.
2. Go to the **Optimize** tab.
3. Under *Partitioning*, choose:
 - o **Current partitioning** (default), or
 - o **Set partitioning** → choose **Hash**, **Range**, or **Key**.
4. Select a column to partition by (e.g., OrderDate, CustomerId).

5. Set number of partitions (or let ADF decide based on compute).
6. Publish and test with debug to compare performance (before vs after).

📌 Example you can say in an interview

“We had a big facts table with ~50M rows. The Dataflow was slow because everything was processed in a single partition. I enabled partitioning in the Optimize tab and partitioned by `OrderDate` month. That allowed Spark to process multiple partitions in parallel and reduced Dataflow execution time by around 35–40%.”

◆ Scenario 22 — Copy activity fails due to timeout or long run time

📌 Explanation (What & Why)

When copying huge data (GBs/TBs), Copy Activity can fail with timeouts, throttling, or slow performance. You must tune performance and sometimes use **staged copy** or **PolyBase/Bulk** to offload work.

📌 How to implement in ADF (Steps)

1. Open **Copy Activity** → **Settings**.
2. Increase **Timeout** (e.g., from default to 2–4 hours for heavy loads).
3. In **Performance**:
 - Enable **Parallel copy** (e.g., 16–32).
 - Set **Data integration units (DIUs)** higher if needed.
4. For SQL → Synapse loads, enable **Staging (PolyBase)**:
 - Use **Blob/ADLS** as staging.
 - Check *Use PolyBase* or *Bulk insert* where available.
5. Ensure network bandwidth / SHIR VM size is sufficient.
6. Re-run and monitor throughput (MB/s) from Monitor tab.

📌 Example interview answer

“Our SQL-to-ADLS copy was failing with timeout due to 200M+ rows. I increased the activity timeout, enabled parallel copy with 16 parallel threads, and used staged copy with PolyBase into Synapse. That increased throughput and eliminated timeouts.”

◆ Scenario 23 — Process near real-time/streaming data with ADF

📌 Explanation (What & Why)

ADF is mainly batch-oriented, but you can achieve *near real-time* by combining **Event Hub** or **IoT Hub** with frequent-trigger pipelines or using Dataflows on streaming landing data.

📌 How to implement in ADF (Steps)

1. Use **Event Hub / IoT Hub** to receive streaming messages.
2. Configure **Capture** to write events into **ADLS/Blob** in small files (e.g., every 1–5 minutes).
3. In ADF, create a **Tumbling Window Trigger** (e.g., every 5 minutes).
4. Pipeline steps:
 - o Read newly landed captured files from ADLS.
 - o Use **Mapping Data Flow** or **Databricks** to clean/transform.
 - o Write results to **Silver/Gold** layer or Synapse.
5. Implement **watermark/last processed time** to avoid double-processing.

📌 Example interview answer

“For IoT telemetry, events landed in Event Hub and were captured to ADLS every 5 minutes. I created a tumbling window trigger in ADF that picked up the latest batch, transformed it in Dataflows, and loaded curated data into Synapse. This gave us near real-time dashboards with a 5-minute delay.”

◆ Scenario 24 — Load CSV files with varying number of columns (schema drift)

📌 Explanation (What & Why)

In real projects, CSV files may evolve — new columns added, some removed. This is called **schema drift**. If you rigidly define columns, pipelines will fail on schema changes. ADF Dataflow can handle schema drift dynamically.

📌 How to implement in ADF (Steps)

1. In **Mapping Data Flow**, create a **Source** pointing to your CSV.
2. In source options, enable **Allow schema drift**.
3. Use **Select or Derived Column** to map only required columns.
4. Use **Column patterns** if you want to handle “all extra columns” generically.
5. Sink:
 - o Use a flexible format like **Parquet** or **Delta** which supports schema evolution.
6. Test with old and new versions of the files.

📌 Example interview answer

“Client CSV files kept changing with extra columns. To avoid frequent failures and rework, I enabled schema drift in Mapping Data Flow and mapped only the required core columns, while still landing all new columns into a ‘flexible’ parquet sink for exploration.”

◆ Scenario 25 — Append data instead of overwriting existing data

📌 Explanation (What & Why)

Sometimes you ingest **daily incremental data** and want to **append** to existing data instead of overwriting the full dataset. This is common for logs, transaction feeds, etc.

📌 How to implement in ADF (Steps)

1. In **Copy Activity** → **Sink settings**:
 - Choose **Copy behavior**: Merge, Upsert, or Insert/Append (options vary by connector).
2. For data lake formats like Parquet:
 - Use file naming strategy with date/time partition (so every run writes new files).
3. For SQL/Synapse:
 - Use **pre-copy script** to filter duplicates if needed.
4. Optionally use **Dataflows** or **MERGE** if you need upsert logic.

📌 Example interview answer

“For daily transaction loads, we just needed to append new data. I configured the Copy Activity sink to append to existing parquet partitions by date, using dynamic paths like `year=/month=/day=`. This allowed us to keep history without overwriting older data.”

◆ Scenario 26 — Automatically delete old files from data lake (retention policy)

📌 Explanation (What & Why)

Storage costs can explode if you keep all raw data forever. Often you need a **retention policy** (e.g., keep only 90 days of raw files).

📌 How to implement in ADF (Steps)

1. Create a pipeline with these activities:
 - **Get Metadata** (list files in a folder).
 - **ForEach** over the list of files.
 - Inside ForEach → **If Condition**:
 - Compare each file’s `lastModified` to `utcNow()` minus X days (e.g., 90).
 - If older → **Delete Activity** to delete that file.
2. Schedule this pipeline with a **daily trigger**.
3. Log deleted files (optional) into a table or log file.

📌 Example interview answer

“To control storage, I built a cleanup pipeline that runs daily. It lists files from the raw folder and deletes anything older than 90 days using Delete Activity. This was fully dynamic and based on file `lastModified` metadata, helping us meet cost and compliance requirements.”

◆ Scenario 27 — Handle NULL / missing values in transformations

📌 Explanation (What & Why)

Nulls can break business logic and cause incorrect aggregates. You often need to **impute** or replace nulls to maintain data quality.

📌 How to implement in ADF (Steps)

1. In **Mapping Data Flow**, add a **Derived Column** transformation.
2. Use expressions like:
 - o `iif(isNull(City), 'Unknown', City)`
 - o `iif(isNull(SalesAmount), 0, SalesAmount)`
3. For numeric columns, you can default to 0 or average; for string columns, use defined placeholders.
4. Optionally add **Assert** transformation to enforce “non-null” constraints for critical columns.

📌 Example interview answer

“Customer master had many null `City` values, which was affecting segmentation. In Dataflows, I added a Derived Column to replace null cities with ‘Unknown’ and null numeric fields with 0. This ensured cleaner aggregates and prevented issues in downstream reports.”

◆ Scenario 28 — Combine thousands of small files into fewer large files

📌 Explanation (What & Why)

Too many small files (“small file problem”) makes queries slow and costly (especially with Spark/Synapse). Best practice is to **compact** small files into larger partitions.

📌 How to implement in ADF (Steps)

1. Use **Mapping Data Flow** or **Databricks** to read all small files from a folder.
2. In **Sink** settings:
 - o Set **File name option** = Single file per partition or pattern.
 - o Configure **Partitioning** so you get fewer, larger files.
3. Write data as **Parquet/Delta** with partitioning by date/other columns.
4. Optionally remove or archive the original small files after successful compaction.

📌 Example interview answer

“We were getting thousands of 1–5 KB files per hour from IoT sensors. Queries were very slow, so I created a daily compaction pipeline that read all small files and wrote them into fewer large parquet files partitioned by date. This reduced the file count drastically and improved query performance.”

◆ Scenario 29 — Process weekly data using Tumbling Window trigger

📌 Explanation (What & Why)

Some business jobs run on **weekly cycles**, like weekly sales reports, aggregates, or snapshots. Tumbling Window triggers are perfect for such fixed, non-overlapping intervals.

📌 How to implement in ADF (Steps)

1. Create a **Tumbling Window Trigger**:
 - o Set **recurrence** to 7 days.
 - o Set **start time** (e.g., every Monday 01:00 AM).
2. In the pipeline, use trigger context functions such as:
 - o `windowStart()` and `windowEnd()` for filtering data.
3. In Copy/Dataflow source query:
 - o Filter data between `windowStart` and `windowEnd`.
4. Run and verify that each window processes only that specific week’s data.

📌 Example interview answer

“We built weekly sales aggregates that had to run every Monday. I used a tumbling window trigger with a 7-day window. Using `windowStart` and `windowEnd`, I filtered data in the source query so each run processed exactly one week of sales, with no overlaps.”

◆ Scenario 30 — Fail pipeline if data volume is below an expected threshold

📌 Explanation (What & Why)

If a feed suddenly drops from millions of rows to a few records, that usually indicates a problem. You can create **data reasonability checks** and fail the pipeline when volume is suspiciously low.

📌 How to implement in ADF (Steps)

1. After a **Copy Activity**, use a **Lookup / Stored Procedure** or **Dataflow** to get **row count** of the loaded data.
2. Compare row count with a configured threshold (e.g., from pipeline parameter or config table).

3. Use **If Condition Activity**:
 - o If RowCount < Threshold → use **Fail Activity** with custom error message.
 - o Else → continue pipeline.
4. Optionally send notification (Logic App / Email) when failure occurs.

📌 Example interview answer

“I implemented a data-quality check where, after loading data, we compared row count with a historical threshold. If the count dropped by more than 80%, an If Condition triggered a Fail activity and sent an alert. This prevented bad/partial loads from silently flowing into production.”

◆ Scenario 31 — Reprocess data for last ‘N’ days (handle late-arriving data)

📌 Explanation

Sometimes data arrives late due to upstream delays (e.g., at night). You need to **reprocess the last 3–7 days** to ensure completeness.

📌 How-To Steps

1. Create a pipeline parameter: `reprocessDays`.
2. Use it in ADF to compute start date:
`@addDays(utcNow(), -pipeline().parameters.reprocessDays)`
3. In your SQL query or Dataflow filter:
`SELECT * FROM Sales`
`WHERE UpdatedOn >= @startDate`
4. Write output to Silver/Gold layers annually.
5. Optionally delete duplicate records using MERGE.

📌 Example for interviewer

“For late-arriving orders, I built a parameterized pipeline that reprocesses the last 3 days using dynamic date filtering. This ensures no missing transactions in downstream analytics.”

◆ Scenario 32 — Handle missing or corrupt files gracefully

📌 Explanation

If a file is corrupt or has zero bytes, ingestion must **fail safely** or **skip** it without breaking the pipeline.

📌 How-To Steps

1. Use **Get Metadata** to read:
 - o size
 - o columnCount
 - o lastModified
2. Add **If Condition**:
 - o If size > 0 → process
 - o Else → move file to /error/ folder
3. Use **Copy Activity** to move corrupted files.
4. Log file details (name, timestamp, error reason).

📌 Example for interviewer

“One day, a zero-byte file caused Copy Activity failures. I added a pre-check using Get Metadata → If size=0 → Move to error folder. This made the pipeline fault-tolerant.”

◆ Scenario 33 — Dynamically generate SQL queries using ADF parameters

📌 Explanation

Some pipelines need dynamic WHERE clauses, table names, or partitions based on input.

📌 How-To Steps

1. Create pipeline parameters:
 - o tableName
 - o startDate
 - o endDate
2. Use dynamic SQL in Copy Activity:
3. @concat(
4. 'SELECT * FROM ', pipeline().parameters.tableName,
5. ' WHERE UpdatedOn >= ''', pipeline().parameters.startDate, ''',
6. ' AND UpdatedOn < ''', pipeline().parameters.endDate, '''
7.)
8. Pass this as a source query.

📌 Example

“We used dynamic content in Copy Activity to extract data for specific tables and date ranges, making the pipeline metadata-driven.”

◆ Scenario 34 — Implement SCD Type 1 (overwrite existing values)

📌 Explanation

SCD1 updates existing records instead of keeping history. Used for non-historical data.

📌 How-To Steps

Using Mapping Dataflow:

1. Source: new data
2. Lookup: existing target data by key
3. Alter Row:
 - When matched → UPDATE
 - When not matched → INSERT

Using Databricks:

```
MERGE INTO dim_customer AS tgt
USING staging_customer AS src
ON tgt.CustomerID = src.CustomerID
WHEN MATCHED THEN UPDATE SET *
WHEN NOT MATCHED THEN INSERT *
```

📌 Example

“Customer phone number updates don’t need history, so we implemented SCD1 using Delta MERGE in Databricks.”

◆ Scenario 35 — Trigger multiple child pipelines in parallel (fan-out)

📌 Explanation

Useful when ingesting multiple tables or files at once.

📌 How-To Steps

1. Lookup activity fetches list of tables/files.
2. Use **ForEach**:
 - Set **Batch count** (e.g., 10).
3. Inside ForEach → add **Execute Pipeline** activity.
4. Each child pipeline processes one table/file.

📌 Example

“We ingested 80 tables by using Lookup + ForEach + Execute Pipeline with batch size 10, improving total load time from 6 hours to 1.5 hours.”

◆ Scenario 36 — Build Metadata-driven ADF pipelines

📌 Explanation

You store configurations in a SQL table instead of hardcoding values in ADF. Makes pipelines reusable across many sources.

📌 How-To Steps

1. Create a config table:
2.

SourceTable	TargetPath	LoadType	WatermarkColumn
-----	-----	-----	-----
4. Sales	/sales/	Incremental	ModifiedDate
5. Customers	/cust/	FullLoad	null
6. Use Lookup to read config row.
7. Pass values into child pipeline parameters.
8. Apply dynamic transforms:
 - dynamic source table
 - dynamic sink folder
 - dynamic load type

📌 Example

“A single metadata-driven pipeline ingests 30 different tables using config stored in SQL. Zero hardcoding.”

◆ Scenario 37 — Auto-create folders in ADLS based on date

📌 Explanation

Folders like `/year=2025/month=01/day=12/` must be created dynamically for partitions.

📌 How-To Steps

Use dynamic expressions:

```
@concat(  
    'year=' , formatDateTime(utcNow() , 'yyyy') ,  
    '/month=' , formatDateTime(utcNow() , 'MM') ,  
    '/day=' , formatDateTime(utcNow() , 'dd')  
)
```

Set this in Sink dataset → folder path.

📌 Example

“Daily partitions were created dynamically in ADLS using folder path expressions like year/month/day.”

◆ Scenario 38 — Call an API that requires authentication token

📌 Explanation

Some APIs require an initial token call → then use the token in subsequent requests.

📌 How-To Steps

1. First Web Activity:
 - o POST to /auth/token
 - o Save token to variable:
 - o @activity('AuthCall').output.access_token
2. Second Web Activity: call actual API:
 - o Add header:
 - o Authorization: Bearer @{variables('token')}
3. Copy Activity loads JSON response to ADLS/SQL.

📌 Example

“To extract data from ServiceNow API, I created a token-generation Web Activity, stored the token, and used it across multiple API calls.”

◆ Scenario 39 — Load nested JSON into SQL (Flatten + Mapping)

📌 Explanation

SQL cannot store nested JSON, so flattening is required.

📌 How-To Steps

1. Use Mapping Dataflow.
2. Add **Flatten** transformation:
 - o Select array/object field.
3. Map flattened columns to SQL table.
4. Use Derived Column to rename nested fields.

📌 Example

Input JSON:

```
{  
  "custId": 101,
```

```

"address": {
  "city": "Delhi",
  "pincode": 110085
}
}

```

Flatten result:

```
custId | address_city | address_pincode
```

◆ Scenario 40 — Stop duplicate files from processing twice

📌 Explanation

If the same file lands twice, your logic must **skip** or **overwrite** correctly to prevent duplicate data.

📌 How-To Steps

1. Maintain a table: ProcessedFileLog
 - o fileName
 - o hash
 - o processedDate
2. Before processing:
 - o Lookup the log table
 - o If file exists → skip
 - o Else → process
3. After successful run → insert entry in log table.

📌 Example

“I prevented reprocessing by storing both filename and MD5 checksum in a log table. If the same file arrives, the pipeline skips it.”

◆ Scenario 41 — Improve Copy Activity performance (10× faster loads)

📌 Explanation

Copy Activity performance depends on:

- Parallelism
- DIUs
- File format
- Network throughput

- Source query efficiency

📌 How-To Steps

1. Increase **DIUs** (Data Integration Units).
2. Set **parallelCopies** to 8–32 based on source.
3. Use **partitioned SQL queries** for large tables.
4. Choose **Parquet** instead of CSV for sink (10× faster).
5. Enable **staged copy** for SQL → Synapse loads.
6. Use **column pruning** (SELECT only required columns).
7. Use **compression** like snappy/gzip.

📌 Example

“We moved 180M rows from SQL to ADLS. By using Parquet, increasing DIUs, and setting 16 parallel copies, throughput increased from 12 MB/s to 110 MB/s.”

◆ Scenario 42 — Use Data Partitioning for faster processing

📌 Explanation

Partitioning divides data into chunks so Spark processes them in parallel, reducing execution time drastically.

📌 How-To Steps

In **Mapping Dataflow** → **Optimize** tab:

1. Choose **Set partitioning**.
2. Select **Hash**, **Key**, or **Range** partitioning.
3. Pick a column:
 - OrderDate
 - CustomerId
4. Optionally define number of partitions.

📌 Example

“Partitioning a 50M row dataset by OrderDate dropped Dataflow runtime from 13 minutes to 5 minutes.”

◆ Scenario 43 — Debug Mapping Dataflows effectively

📌 Explanation

Debugging helps validate transformations and preview data.

📌 How-To Steps

1. Enable **Debug Mode**.
2. Use **Data Preview** at each transformation.
3. Use **Inspect** tab to review schema and data types.
4. Use breakpoints (temporary endpoints).
5. Validate expressions using built-in expression tester.

📌 Example

“I discovered a null handling error in Derived Column by turning on Data Preview and inspecting mismatched values.”

◆ Scenario 44 — Scale out Integration Runtime automatically

📌 Explanation

Azure IR automatically scales based on workload, but SHIR requires manual scaling.

📌 How-To Steps

For Azure IR (Auto-scale):

- No action required; Microsoft manages compute.
- You only adjust **DIUs** during copy.

For Self-Hosted IR:

1. Add extra SHIR nodes.
2. Enable **High Availability** mode.
3. Enable **Auto Update**.
4. Distribute load by mapping multiple pipelines to SHIR nodes.

📌 Example

“To handle heavy on-prem Oracle loads, I added two more SHIR nodes, doubling throughput.”

◆ Scenario 45 — PolyBase vs Bulk Copy for Synapse ingestion

📌 Explanation

Both are used to ingest data into Synapse, but differ in performance and function.

📌 How-To Steps

Feature	PolyBase	Bulk Copy
Performance	Extremely fast for large datasets	Medium-fast
Use case	SQL → Synapse loads	CSV → Synapse loads
Requires staging?	Yes	No
Parallelism	Very high	Limited

In ADF:

- Enable **Use PolyBase** when copying into Synapse from Blob/ADLS.
- Use **Bulk Insert** for smaller but frequent loads.

📌 Example

“PolyBase reduced load time for 40 GB parquet data from 1 hour to 9 minutes.”

◆ Scenario 46 — Use Managed Identity to secure ADF access

📌 Explanation

Managed Identity eliminates passwords and securely authenticates ADF to resources like ADLS, Key Vault, SQL, etc.

📌 How-To Steps

1. Enable **Managed Identity** in ADF.
2. Assign roles to ADF identity:
 - **Storage Blob Contributor**
 - **Key Vault Reader**
 - **SQL DB Contributor**
3. In linked services, choose **Managed Identity authentication**.
4. Remove hardcoded credentials.

📌 Example

“We replaced SQL credentials with Managed Identity, improving security and meeting compliance.”

◆ Scenario 47 — Setup RBAC for Data Factory (Granular permissions)

📌 Explanation

RBAC controls who can edit pipelines, run pipelines, read data, etc.

📌 How-To Steps

Go to **Azure Portal** → **ADF** → **Access Control (IAM)**

Assign roles:

- **Owner** — full access
- **Contributor** — create/update pipelines
- **Reader** — view only
- **Data Factory Pipeline Operator** — can run pipelines, not edit

📌 Example

“We restricted junior developers to Pipeline Operator role so they couldn’t modify production pipelines.”

◆ Scenario 48 — Centralized monitoring using ADF + Azure Monitor

📌 Explanation

ADF integrates logs & metrics into Azure Monitor for centralized alerting.

📌 How-To Steps

1. Enable **Diagnostic Settings** in ADF:
 - PipelineRuns
 - ActivityRuns
 - TriggerRuns
2. Send logs to:
 - Log Analytics Workspace
 - Event Hub
 - Storage account
3. Build **KQL dashboards** for failures or long-running pipelines.

📌 Example

“We created KQL alerts for pipeline failures. Teams received alerts within 30 seconds of failure.”

◆ Scenario 49 — Implement audit logging for every activity in pipeline

📌 Explanation

Logging ensures traceability and governance.

📌 How-To Steps

1. Use **Activity Output** to extract:
 - runId
 - activityName
 - status
 - start/end time
2. Insert logs into SQL or ADLS via:
 - Stored Procedure
 - Web activity (if hitting API)
 - Copy to lake
3. Use **OnSuccess**, **OnFailure**, and **OnCompletion** paths.

📌 Example

“Each activity wrote log entries to a central audit table. This helped debug failures quickly.”

◆ Scenario 50 — Encrypt data in transit and at rest

📌 Explanation

ADF supports encryption automatically, but you may need to enable specific features for compliance.

📌 How-To Steps

At rest:

- ADLS, Blob, SQL, Synapse → automatically encrypted with AES-256.
- For extra safety, enable **Customer Managed Keys (CMK)**.

In transit:

- ADF enforces **HTTPS/TLS 1.2**.
- Use Private Endpoints to avoid public internet.
- For SQL, enforce TLS connection strings.

📌 Example

“For a finance client, we enabled Customer Managed Keys for ADLS encryption and private endpoints for secure pipeline traffic.”